

Force-Directed Drawing Algorithms

	5.1 Introduction.....	1
	5.2 Spring Systems and Electrical Forces	2
	5.3 The Barycentric Method	4
	5.4 Further Spring Refinements.....	5
	5.5 Graph Theoretic Distances Approach	6
Stephen G. Kobourov	5.6 Large Graphs	7
University of Arizona	5.7 Non-Euclidean Approaches	11

5.1 Introduction

Some of the most flexible algorithms for calculating layouts of simple undirected graphs belong to a class known as force-directed algorithms. Also known as spring embedders, such algorithms calculate the layout of a graph using only information contained within the structure of the graph itself, rather than relying on domain-specific knowledge. Graphs drawn with these algorithms tend to be aesthetically pleasing, exhibit symmetries, and tend to produce crossing-free layouts for planar graphs. In this chapter we will assume that the input graphs are simple, connected, undirected graphs and their layouts are straight-line drawings. Excellent surveys of this topic can be found in Di Battista *et al.* [DETT99] Chapter 10 and Brandes [Bra01].

Going back to 1963, the graph drawing algorithm of Tutte [Tut63] is one of the first force-directed graph drawing methods based on barycentric representations. More traditionally, the spring layout method of Eades [Ead84] and the algorithm of Fruchterman and Reingold [FR91] both rely on spring forces, similar to those in Hooke's law. In these methods, there are repulsive forces between all nodes, but also attractive forces between nodes which are adjacent.

Alternatively, forces between the nodes can be computed based on their graph theoretic distances, determined by the lengths of shortest paths between them. The algorithm of Kamada and Kawai [KK89] uses spring forces proportional to the graph theoretic distances. In general, force-directed methods define an objective function which maps each graph layout into a number in \mathcal{R}^+ representing the energy of the layout. This function is defined in such a way that low energies correspond to layouts in which adjacent nodes are near some pre-specified distance from each other, and in which non-adjacent nodes are well-spaced. A layout for a graph is then calculated by finding a (often local) minimum of this objective function.

The utility of the basic force-directed approach is limited to small graphs and results are poor for graphs with more than a few hundred vertices. There are multiple reasons why traditional force-directed algorithms do not perform well for large graphs. One of the main

obstacles to the scalability of these approaches is the fact that the physical model typically has many local minima. Even with the help of sophisticated mechanisms for avoiding local minima the basic force-directed algorithms are not able to consistently produce good layouts for large graphs. Barycentric methods also do not perform well for large graphs mainly due to resolution problems: for large graphs the minimum vertex separation tends to be very small, leading to unreadable drawings.

The late 1990's saw the emergence of several techniques extending the functionality of force-directed methods to graphs with tens of thousands and even hundreds of thousands of vertices. One common thread in these approaches is the multi-level layout technique, where the graph is represented by a series of progressively simpler structures and laid out in reverse order: from the simplest to the most complex. These structures can be coarser graphs (as in the approach of Hadany and Harel [HH01], Harel and Koren [HK02], and Walshaw [Wal03], or vertex filtrations as in the approach of Gajer, Goodrich, and Kobourov [GGK04].

The classical force-directed algorithms are restricted to calculating a graph layout in Euclidean geometry, typically \mathcal{R}^2 , \mathcal{R}^3 , and, more recently, \mathcal{R}^n for larger values of n . There are, however, cases where Euclidean geometry may not be the best option: Certain graphs may be known to have a structure which would be best realized in a different geometry, such as on the surface of a sphere or on a torus. In particular, 3D mesh data can be parameterized on the sphere for texture mapping or graphs of genus one can be embedded on a torus without crossings. Furthermore, it has also been noted that certain non-Euclidean geometries, specifically hyperbolic geometry, have properties which are particularly well suited to the layout and visualization of large classes of graphs [LRP95, Mun97]. With this in mind, Kobourov and Wampler describe extensions of the force-directed algorithms to Riemannian spaces [KW05].

5.2 Spring Systems and Electrical Forces

The 1984 algorithm of Eades [Ead84] targets graphs with up to 30 vertices and uses a mechanical model to produce “aesthetically pleasing” 2D layouts for plotters and CRT screens. The algorithm is succinctly summarized as follows:

To embed a graph we replace the vertices by steel rings and replace each edge with a spring to form a mechanical system. The vertices are placed in some initial layout and let go so that the spring forces on the rings move the system to a minimal energy state. Two practical adjustments are made to this idea: firstly, logarithmic strength springs are used; that is, the force exerted by a spring is:

$$c_1 * \log(d/c_2),$$

where d is the length of the spring, and c_1 and c_2 are constants. Experience shows that Hookes Law (linear) springs are too strong when the vertices are far apart; the logarithmic force solves this problem. Note that the springs exert no force when $d = c_2$. Secondly, we make nonadjacent vertices repel each other. An inverse square law force,

$$c_3/\sqrt{d},$$

where c_3 is constant and d is the distance between the vertices, is suitable. The mechanical system is simulated by the following algorithm.

```

algorithm SPRING( $G$ :graph);
place vertices of  $G$  in random locations;
repeat  $M$  times
    calculate the force on each vertex;
    move the vertex  $c_4 * (\text{force on vertex})$ 
draw graph on CRT or plotter.

```

The values $c_1 = 2$, $c_2 = 1$, $c_3 = 1$, $c_4 = 0.1$, are appropriate for most graphs. Almost all graphs achieve a minimal energy state after the simulation step is run 100 times, that is, $M = 100$.

This excellent description encapsulates the essence of spring algorithms and their natural simplicity, elegance, and conceptual intuitiveness. The goals behind “aesthetically pleasing” layouts were initially captured by the two criteria: “all the edge lengths ought to be the same, and the layout should display as much symmetry as possible.”

The 1991 algorithm of Fruchterman and Reingold added “even vertex distribution” to the earlier two criteria and treats vertices in the graph as “atomic particles or celestial bodies, exerting attractive and repulsive forces from one another.” The attractive and repulsive forces are redefined to

$$f_a(d) = d^2/k, \qquad f_r(d) = -k^2/d,$$

in terms of the distance d between two vertices and the optimal distance between vertices k defined as

$$k = C \sqrt{\frac{\text{area}}{\text{number of vertices}}}.$$

This algorithm is similar to that of Eades in that both algorithms compute attractive forces between adjacent vertices and repulsive forces between all pairs of vertices. The algorithm of Fruchterman and Reingold adds the notion of “temperature” which could be used as follows: “the temperature could start at an initial value (say one tenth the width of the frame) and decay to 0 in an inverse linear fashion.” The temperature controls the displacement of vertices so that as the layout becomes better, the adjustments become smaller. The use of temperature here is a special case of a general technique called simulated annealing, whose use in force-directed algorithms is discussed later in this chapter. The pseudo-code for the Fruchterman and Reingold provides further insight into the workings of a spring-embedder:

```

area :=  $W * L$ ; { $W$  and  $L$  are the width and length of the frame}
 $G := (V, E)$ ; {the vertices are assigned random initial positions}
 $k := \sqrt{\text{area}/|V|}$ ;
function  $f_a(x) := \text{begin return } x^2/k \text{ end}$ ;
function  $f_r(x) := \text{begin return } k^2/x \text{ end}$ ;
for  $i := 1$  to  $\text{iterations}$  do begin
    {calculate repulsive forces}
    for  $v$  in  $V$  do begin
        {each vertex has two vectors: .pos and .disp}
         $v.\text{disp} := 0$ ;
        for  $u$  in  $V$  do
            if  $(u \neq v)$  then begin
                { $\delta$  is the difference vector between the positions of the two vertices }
                 $\delta := v.\text{pos} - u.\text{pos}$ ;
                 $v.\text{disp} := v.\text{disp} + (\delta/|\delta|) * f_r(|\delta|)$ 
            end
        end
    end
    {calculate attractive forces}
    for  $e$  in  $E$  do begin
        {each edges is an ordered pair of vertices .v and .u}
         $\delta := e.v.\text{pos} - e.u.\text{pos}$ ;
         $e.v.\text{disp} := e.v.\text{disp} - (\delta/|\delta|) * f_a(|\delta|)$ ;
         $e.u.\text{disp} := e.u.\text{disp} + (\delta/|\delta|) * f_a(|\delta|)$ 
    end
    {limit max displacement to temperature  $t$  and prevent from displacement outside frame}
    for  $v$  in  $V$  do begin
         $v.\text{pos} := v.\text{pos} + (v.\text{disp}/|v.\text{disp}|) * \min(v.\text{disp}, t)$ ;
         $v.\text{pos}.x := \min(W/2, \max(-W/2, v.\text{pos}.x))$ ;
         $v.\text{pos}.y := \min(L/2, \max(-L/2, v.\text{pos}.y))$ 
    end
    {reduce the temperature as the layout approaches a better configuration}
     $t := \text{cool}(t)$ 
end

```

Each iteration the basic algorithm computes $O(|E|)$ attractive forces and $O(|V|^2)$ repulsive forces. To reduce the quadratic complexity of the repulsive forces, Fruchterman and Reingold suggest using a grid variant of their basic algorithm, where the repulsive forces between distant vertices are ignored. For sparse graphs, and with uniform distribution of the vertices, this method allows a $O(|V|)$ time approximation to the repulsive forces calculation. This approach can be thought of as a special case of the multi-pole technique introduced in n -body simulations [Gre88] whose use in force-directed algorithms will be further discussed later in this chapter.

As in the paper by Eades [Ead84] the graphs considered by Fruchterman and Reingold are small graphs with less than 40 vertices. The number of iterations of the main loop is also similar at 50.

5.3 The Barycentric Method

Historically, Tutte's 1963 barycentric method [Tut63] is the first "force-directed" algorithm for obtaining a straight-line, crossings free drawing for a given 3-connected planar graph. Unlike almost all other force-directed methods, Tutte's guarantees that the resulting drawing is crossings-free; moreover, all faces of the drawing are convex.

The idea behind Tutte's algorithm is that if a face of the planar graph is fixed in the plane, then suitable positions for the remaining vertices can be found by solving a system of

linear equations, where each vertex position is represented as a convex combination of the positions of its neighbors. This can be considered a force-directed method as summarized in Di Battista *et al.* [DETT99]:

Barycenter-Draw

Input: $G = (V, E)$; a partition $V = V_0 \sqcup V_1$ of V into a set V_0 of at least three *fixed* vertices and a set V_1 of *free* vertices; a strictly convex polygon P with $|V_0|$ vertices

Output: a position p_v for each vertex of V , such that the fixed vertices form a convex polygon P .

1. Place each fixed vertex $u \in V_0$ at a vertex of P , and each free vertex at the origin.

2. **repeat**

foreach free vertex $v \in V_1$ **do**

$$x_v = \frac{1}{\deg(v)} \sum_{(u,v) \in E} x_u$$

$$y_v = \frac{1}{\deg(v)} \sum_{(u,v) \in E} y_u$$

until x_v and y_v converge for all free vertices v .

In this model the force due to an edge (u, v) is proportional to the distance between vertices u and v and the springs have ideal length of zero; there are no explicit repulsive forces. Thus the force at a vertex v is described by

$$F(v) = \sum_{(u,v) \in E} (p_u - p_v),$$

where p_u and p_v are the positions of vertices u and v . As this function has a trivial minimum with all vertices placed in the same location, the vertex set is partitioned into fixed and free vertices. Setting the partial derivatives of the force function to zero results in independent systems of linear equations for the x -coordinate and for the y -coordinate.

The equations in the for-loop are linear and the number of equations is equal to the number of the unknowns, which in turn is equal to the number of free vertices. Solving these equations results in placing each free vertex at the barycenter of its neighbors. The system of equations can be solved using the Newton-Raphson method. Moreover, the resulting solution is unique.

One significant drawback of this approach is the resulting drawing often has poor vertex resolution. In fact, for every $n > 1$, there exists a graph, such that the barycenter method computes a drawing with exponential area [EG95].

5.4 Further Spring Refinements

Even before the 1984 algorithm of Eades, force-directed techniques were used in the context of VLSI layouts in the 1960's and 1970's [FCW67, QB79]. Yet, renewed interest in force-directed graph layout algorithms brought forth many new ideas in the 1990's. Frick, Ludwig, and Mehldau [FLM95] add new heuristics to the Fruchterman-Reingold approach. In particular, oscillation and rotations are detected and dealt with using local instead of global temperature. The following year Bruß and Frick [BF96] extended the approach to layouts directly in 3D Euclidean space. The algorithm of Cohen [Coh97] introduced the notion of an incremental layout, a precursor of the multi-scale methods described in section 5.6.

The 1997 algorithm of Davidson and Harel [DH96] adds additional constraints to the traditional force-directed approach in explicitly aiming to minimize the number of edge-crossings and keeping vertices from getting too close to non-adjacent edges. The algorithm uses the simulated annealing technique developed for large combinatorial optimization [KGV83]. Simulated annealing is motivated by the physical process of cooling molten materials. When molten steel is cooled too quickly it cracks and forms bubbles making it brittle. For better results, the steel must be cooled slowly and evenly and this process is known as annealing in metallurgy. With regard to force-directed algorithms, this process is simulated to find local minima of the energy function.

Genetic algorithms for force-directed placement have also been considered. Genetic algorithms are a commonly used search technique for finding approximate solutions to optimization and search problems. The technique is inspired by evolutionary biology in general and by inheritance, mutation, natural selection, and recombination (or crossover), in particular; see the survey by Vose [v-s99]. In the context of force-directed techniques for graph drawing, the genetic algorithms approach was introduced in 1991 by Kosak, Marks and Shieber [KMS91]. Other notable approaches in the direction include that of Branke, Bucher, and Schmeck [BBS97].

5.5 Graph Theoretic Distances Approach

The 1989 algorithm of Kamada and Kawai [KK89] introduced a different way of thinking about “good” graph layouts. Whereas the algorithms of Eades and Fruchterman-Reingold aim to keep adjacent vertices close to each other while ensuring that vertices are not too close to each other, Kamada and Kawai take graph theoretic approach:

“We regard the desirable geometric (Euclidean) distance between two vertices in the drawing as the “graph theoretic distance between them in the corresponding graph.”

In this model, the “perfect” drawing of a graph would be one in which the pair-wise geometric distances between the drawn vertices match the graph theoretic pairwise distances, as computed by an All-Pairs-Shortest-Path computation. As this goal cannot always be achieved for arbitrary graphs in 2D or 3D Euclidean spaces, the approach relies on setting up a spring system in such a way that minimizing the energy of the system corresponds to minimizing the difference between the geometric and graph distances. In this model there are no separate attractive and repulsive forces between pairs of vertices, but instead if a pair of vertices is (geometrically) closer/farther than their corresponding graph distance the vertices repel/attract each other. Let $d_{i,j}$ denote the shortest path distance between vertex i and vertex j in the graph. Then $l_{i,j} = L \times d_{i,j}$ is the ideal length of a spring between vertices i and j , where L is the desirable length of a single edge in the display. Kamada and Kawai suggest that $L = L_0 / \max_{i < j} d_{i,j}$, where L_0 is the length of a side of the display area and $\max_{i < j} d_{i,j}$ is the diameter of the graph, i.e., the distance between the farthest pair of vertices. The strength of the spring between vertices i and j is defined as

$$k_{i,j} = K/d_{i,j}^2,$$

where K is a constant. Treating the drawing problem as localizing $|V| = n$ particles p_1, p_2, \dots, p_n in 2D Euclidean space, leads to the following overall energy function:

$$E = \sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{1}{2} k_{i,j} (|p_i - p_j| - l_{i,j})^2.$$

The coordinates of a particle p_i in the 2D Euclidean plane are given by x_i and y_i which allows us to rewrite the energy function as follows:

$$E = \sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{1}{2} k_{i,j} \left((x_i - x_j)^2 + (y_i - y_j)^2 + l_{i,j}^2 - 2l_{i,j} \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2} \right).$$

The goal of the algorithm is to find values for the variables that minimize the energy function $E(x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_n)$. In particular, at a local minimum all the partial derivatives are equal to zero, and which corresponds to solving $2n$ simultaneous non-linear equations. Therefore, Kamada and Kawai compute a stable position one particle p_m at a time. Viewing E as a function of only x_m and y_m a minimum of E can be computed using the Newton-Raphson method. At each step of the algorithm the particle p_m with the largest value of Δ_m is chosen, where

$$\Delta_m = \sqrt{\left(\frac{\partial E}{\partial x_m}\right)^2 + \left(\frac{\partial E}{\partial y_m}\right)^2}.$$

This leads to the algorithm below:

```

compute  $d_{i,j}$  for  $1 \leq i \neq j \leq n$ ;
compute  $l_{i,j}$  for  $1 \leq i \neq j \leq n$ ;
compute  $k_{i,j}$  for  $1 \leq i \neq j \leq n$ ;
initialize  $p_1, p_2, \dots, p_n$ ;
while ( $\max_i \Delta_i > \epsilon$ ) {
  let  $p_m$  be the particle satisfying  $\Delta_m = \max_i \Delta_i$ ;
  while ( $\Delta_m > \epsilon$ ) {
    compute  $\delta x$  and  $\delta y$  by solving the following system of equations:
       $\frac{\partial^2 E}{\partial x_m^2}(x_m^{(t)}, y_m^{(t)})\delta x + \frac{\partial^2 E}{\partial x_m \partial y_m}(x_m^{(t)}, y_m^{(t)})\delta y = -\frac{\partial E}{\partial x_m}(x_m^{(t)}, y_m^{(t)});$ 
       $\frac{\partial^2 E}{\partial y_m \partial x_m}(x_m^{(t)}, y_m^{(t)})\delta x + \frac{\partial^2 E}{\partial y_m^2}(x_m^{(t)}, y_m^{(t)})\delta y = -\frac{\partial E}{\partial y_m}(x_m^{(t)}, y_m^{(t)});$ 
       $x_m := x_m + \delta x;$ 
       $y_m := y_m + \delta y;$ 
  }
}
```

The algorithm of Kamada and Kawai is computationally expensive, requiring an All-Pair-Shortest-Path computation which can be done in $O(|V|^3)$ time using the Floyd-Warshall algorithm or in $O(|E||V|)$ via repeated Breadth-First Search; see the All-Pairs-Shortest-Path chapter in an algorithms textbook such as [CLRS90]. Furthermore, the algorithm requires $O(|V|^2)$ storage for the pairwise vertex distances. Despite the higher time and space complexity, the algorithm contributes a simple and intuitive definition of a “good” graph layout: A graph layout is good if the geometric distances between vertices closely correspond to the underlying graph distances.

5.6 Large Graphs

The first force-directed algorithms to produce good layouts for graphs with over 1000 vertices is the 1999 algorithm of Hadany and Harel [HH01]. They introduced the multi-scale technique as a way to deal with large graphs and in the following year four related but independent force-directed algorithms for large graphs were presented at the Annual Symposium on Graph Drawing. We begin with Hadany and Harel’s description on the multi-scale method:

A natural strategy for drawing a graph nicely is to first consider an abstraction, disregarding some of the graph's fine details. This abstraction is then drawn, yielding a “rough” layout in which only the general structure is revealed. Then the details are added and the layout is corrected. To employ such a strategy is it crucial that the abstraction retain essential features of the graph. Thus, one has to define the notion of coarse-scale representations of a graph, in which the combinatorial structure is significantly simplified but features important for visualization are well preserved. The drawing process will then “travel” between these representations, and introduce multi-scale corrections. Assuming we have already defined the multiple levels of coarsening, the general structure of our strategy is as follows:

- 1. Perform fine-scale relocations of vertices that yield a locally organized configuration.*
- 2. Perform coarse-scale relocations (through local relocations in the coarse representations), correcting global disorders not found in stage 1.*
- 3. Perform fine-scale relocations that correct local disorders introduced by stage 2.*

Hadany and Harel suggest computing the sequence of graphs by using edge contractions so as to preserve certain properties of the graph. In particular, the goal is to preserve three topological properties: cluster size, vertex degrees, and homotopy. For the coarse-scale relocations, the energy function for each graph in the sequence is that of Kamada and Kawai (the pairwise graph distances are compared to the geometric distances in the current layout). For the fine-scale relocations, the authors suggest using force-directed calculations as those of Eades [Ead84], Fruchterman-Reingold [FR91], or Kamada-Kawai [KK89]. While the asymptotic complexity of this algorithm is similar to that of the Kamada-Kawai algorithm, the multi-scale approach leads to good layouts for much larger graphs in reasonable time.

The 2000 algorithm of Harel and Koren [HK02] took force-directed algorithms to graphs with 15,000 vertices. This algorithm is similar to the algorithm of Hadany and Harel, yet uses a simpler coarsening process based on a k -centers approximation, and a faster fine-scale beautification. Given a graph $G = (V, E)$, the k -centers problem asks to find a subset of the vertex set $V' \subseteq V$ of size k , so as to minimize the maximum distance from a vertex to V' : $\min_{V'} \max_{u \in V} \max_{v \in V'} \text{dist}(u, v)$. While k -centers is an NP-hard problem, Harel and Koren use a straightforward and efficient 2-approximation algorithm that relies on Breadth-First Search [Hoc96]. The fine-scale vertex relocations are done using the Kamada-Kawai approach. The Harel and Koren algorithm is summarized in pseudo-code below:

```

Layout( $G(V, E)$ )
% Goal: Find  $L$ , a nice layout of  $G$ 
% Constants:
% Rad[= 7] – determines radius of local neighborhoods
% Iterations[= 4] – determines number of iterations in local beautification
% Ratio[= 3] – ratio between number of vertices in two consecutive levels
% MinSize[= 10] – size of the coarsest graph
  Compute the all-pairs shortest path length:  $d_{VV}$ 
  Set up a random layout  $L$ 
   $k \leftarrow \text{MinSize}$ 
  while  $k \leq |V|$  do
     $\text{centers} \leftarrow \mathbf{K}\text{-Centers}(G(V, E), k)$ 
     $\text{radius} = \max_{v \in \text{centers}} \min_{u \in \text{centers}} \{d_{vu}\} * \text{Rad}$ 
    LocalLayout( $d_{\text{centers} \times \text{centers}}, L(\text{centers}), \text{radius}, \text{Iterations}$ )
    for every  $v \in V$  do
       $L(v) \in L(\text{center}(v)) + \text{rand}$ 
     $k \leftarrow k \text{Ratio}$ 
  return  $L$ 

K-Centers( $G(V, E), k$ )
% Goal: Find a set  $S \subseteq V$  of size  $k$ , such that  $\max_{v \in V} \min_{s \in S} \{d_{sv}\}$  is minimized.
   $S \leftarrow \{v\}$  for some arbitrary  $v \in V$ 
  for  $i = 2$  to  $k$  do
    1. Find the vertex  $u$  farthest away from  $S$ 
      (i.e., such that  $\min_{s \in S} \{d_{us}\} \geq \min_{s \in S} \{d_{ws}\}, \forall w \in V$ )
    2.  $S \leftarrow S \cup \{u\}$ 
  return  $S$ 

LocalLayout( $d_{V \times V}, L, k, \text{Iterations}$ )
% Goal: Find a locally nice layout  $L$  by beautifying  $k$ -neighborhoods
%  $d_{V \times V}$  : all-pairs shortest path length
%  $L$ : initialized layout
%  $k$ : radius of neighborhoods
  for  $i = 1$  to  $\text{Iterations} * |V|$  do
    1. Choose the vertex  $v$  with the maximal  $\Delta_v^k$ 
    2. Compute  $\delta_v^k$  as in Kamada-Kawai
    3.  $L(v) \leftarrow L(v) + (\delta_v^k(x), \delta_v^k(y))$ 
  end

```

The 2000 algorithm of Gajer, Goodrich, and Kobourov [GGK04] is also a multi-scale force-directed algorithm but introduces several ideas to the realm of multi-scale force directed algorithms for large graphs. Most importantly, this approach avoids the quadratic space and time complexity of previous force-directed approaches with the help of a simpler coarsening strategy. Instead of computing a series of coarser graphs from the given large graph $G = (V, E)$, Gajer *et al.* produce a vertex filtration $\mathcal{V} : V_0 \supset V_1 \supset \dots \supset V_k \supset \emptyset$, where $V_0 = V(G)$ is the original vertex set of the given graph G . By restricting the number of vertices considered in relocating any particular vertex in the filtration and ensuring that the filtration has $O(\log |V|)$ levels an overall running time of $O(|V| \log^2 |V|)$ is achieved. Filtrations based on graph centers (as in Harel and Koren []) and maximal independent sets are considered. $V = V_0 \supset V_1 \supset \dots \supset V_k \supset \emptyset$, is a maximal independent set filtration of G if V_i is a maximal subset of V_{i-1} for which the graph distance between any pair of its elements is greater than or equal to 2^i .

```

MAIN ALGORITHM
create a filtration  $\mathcal{V} : V_0 \supset V_1 \supset \dots \supset V_k \supset \emptyset$ 
for  $i = k$  to 0 do
  for each  $v \in V_i - V_{i+1}$  do
    find vertex neighborhood  $N_i(v), N_{i-1}(v), \dots, N_0(v)$ 
    find initial position  $pos[v]$  of  $v$ 
  repeat rounds times
    for each  $v \in V_i$  do
      compute local temperature  $heat[v]$ 
       $disp[v] \leftarrow heat[v] \cdot \vec{F}_{N_i}(v)$ 
    for each  $v \in V_i$  do
       $pos[v] \leftarrow pos[v] + disp[v]$ 
  add all edges  $e \in E$ 

```

In addition to the filtration and neighborhood calculations, Gajer *et al.* introduce the idea of realizing the graph in high-dimensional Euclidean space and obtaining 2D or 3D projections at end. The algorithm also relies on intelligent initial placement of vertices based on graph theoretic distances, rather than random initial placement. Finally, the notion of cooling is re-introduced in the context of multi-scale force-directed algorithms.

Another 2000 multilevel algorithm is that of Walshaw [Wal03]. Instead of relying on the Kamada-Kawai type force interactions, this algorithm extends the grid variant of Fruchterman-Reingold to a multilevel algorithm. The coarsening step is based on repeatedly collapsing maximally independent sets of edges, and the fine-scale refinements are based on Fruchterman-Reingold force calculations. The $O(|V|^2)$ algorithm is summarized below:

```

function  $f_g(x, w) :=$  begin return  $-Cwk^2/x$  end
function  $f_i(x, d, w) :=$  begin return  $\{(x - k)/d\} - f_g(x, w)$  end
 $t := t_0;$ 
 $Posn := NewPosn;$ 
while ( $converged \neq 1$ ) begin
   $converged := 1;$ 
  for  $v \in V$  begin
     $OldPosn[v] = NewPosn[v]$ 
  end
  for  $v \in V$  begin
    {initialize  $D$ , the vector of displacements of  $v$ }
     $D := 0;$ 
    {calculate global (repulsive) forces}
    for  $u \in V, u \neq v$  begin
       $\Delta := Posn[u] - Posn[v];$ 
       $D := D + (\Delta/|Delta|) * f_g(|\Delta|, |u|);$ 
    end
    {calculate local (spring) forces }
    for  $u \in \Gamma(v)$  begin
       $\Delta := Posn[u] - Posn[v];$ 
       $D := D + (\Delta/|Delta|) * f_i(|\Delta|, |\Gamma(v)|, |u|);$ 
    end
    {reposition  $v$ }
     $NewPosn[v] = NewPosn[v] + (D/|D|) * \min(t, |D|);$ 
     $\Delta := NewPosn[v] - OldPosn[v];$ 
    if ( $|\Delta| > k \times tol$ )  $converged := 0;$ 
  end
  {reduce the temperature to reduce the maximum movement}
   $t := cool(t);$ 
end

```

The fourth 2000 multilevel force-directed algorithm is due to Quigley and Eades [QE00]. This algorithm relies on the Barnes-Hut n -body simulation method [BH86] and reduces repulsive force calculations to $O(|V| \log |V|)$ time instead of the usual $O(|V|^2)$. One possible drawback to this approach is that the running time depends on the distribution of the vertices. Hachul and Jünger [HJ04] address this problem in their 2004 multilevel algorithm.

5.7 Non-Euclidean Approaches

Much of the work on non-Euclidean graph drawing has been done in hyperbolic space which offers certain advantages over Euclidean space; see Munzner [Mun97, MB96]. For example, in hyperbolic space it is possible to compute a layout for a complete tree with both uniform edge lengths and uniform distribution of nodes. Furthermore, some of the embeddings of hyperbolic space into Euclidean space naturally provide a fish-eye view of the space, which is useful for “focus+context” visualization, as shown by Lamping *et al.* [LRP95]. The early approaches for calculating the layouts of graphs in hyperbolic space, however, are either restricted by their nature to the layout of trees and tree-like graphs, or to layouts on a lattice.

The hyperbolic tree layout algorithms function on the principle of hyperbolic sphere packing, and operate by making each node of a tree, starting with the root, the center of a sphere in hyperbolic space. The children of this node are then given positions on the surface of this sphere and the process recurses on these children. By carefully computing the radii of these spheres it is possible to create aesthetically pleasing layouts for the given tree.

Although some applications calculate the layout of a general graph using this method, the layout is calculated using a spanning tree of the graph and the extra edges are then added in without altering the layout [Mun98]. This method works well for tree-like and quasi-hierarchical graphs, or for graphs where domain-specific knowledge provides a way to create a meaningful spanning tree. However, for general graphs (e.g., bipartite or densely connected graphs) and without relying on domain specific knowledge, the tree-based approach may result in poor layouts.

Methods for generalizing Euclidean geometric algorithms to hyperbolic space, although not directly related to graph drawing, have also been studied. Recently, van Wijk and Nuij [vWN04] proposed a Poincaré’s half-plane projection to define a model for 2D viewing and navigation. Eppstein [Epp03] shows that many algorithms which operate in Euclidean space can be extended to hyperbolic space by exploiting the properties of a Euclidean model of the space (such as the Beltrami-Klein or Poincaré).

Hyperbolic and spherical space have also been used to display self-organizing maps in the context of data visualization. Ontrup and Ritter [OR01] and Ritter [Rit99] extend the traditional use of a regular (Euclidean) grid, on which the self-organizing map is created, with a tessellation in spherical or hyperbolic space. An iterative process is then used to adjust which elements in the data-set are represented by the intersections. Although the hyperbolic space method seems a promising way to display high-dimensional data-sets, the restriction to a lattice is often undesirable for graph visualization.

Ostry [Ost96] considers constraining force-directed algorithms to the surface of three-dimensional objects. This work is based on a differential equation formulation of the motion of the nodes in the graph, and is flexible in that it allows a layout on almost any object, even multiple objects. Since the force calculations are made in Euclidean space, however, this method is inapplicable to certain geometries (e.g., hyperbolic geometry).

Another example of graph embedding within a non-Euclidean geometry is described in the context of generating spherical parameterizations of 3D meshes. Gotsman *et al.* [GGS03]

describe a method for producing such an embedding using a generalization to spherical space of planar methods for expressing convex combinations of points. The implementation of the procedure is similar to the method described in this paper, but it may not lend itself to geometries other than spherical.

Kobourov and Wampler [KW05] describe a conceptually simple approach to generalizing force-directed methods for graph layout from Euclidean geometry to Riemannian geometries. Unlike previous work on non-Euclidean force-directed methods, this approach is not limited to special classes of graphs but can be applied to arbitrary graphs. The method relies on extending the Euclidean notions of distance, angle, and force-interactions to smooth non-Euclidean geometries via projections to and from appropriately chosen tangent spaces. Formal description of the calculations needed to extend such algorithms to hyperbolic and spherical geometries are also detailed.

In 1894 Riemann described a generalization of the geometry of surfaces, which had been studied earlier by Gauss, Bolyai, and Lobachevsky. Two well-known special cases of Riemannian geometries are the two standard non-Euclidean types, spherical geometry and hyperbolic geometry. This generalization led to the modern concept of a Riemannian manifold. Riemannian geometries have less convenient structure than Euclidean geometry, but they do retain many of the characteristics which are useful for force-directed graph layouts. A Riemannian manifold M has the property that for every point $x \in M$, the tangent space $T_x M$ is an inner product space. This means that for every point on the manifold, it is possible to define local notions of length and angle.

Using the local notions of length we can define the length of a continuous curve $\gamma : [a, b] \rightarrow M$ by

$$\text{length}(\gamma) = \int_a^b \|\gamma'(t)\| dt.$$

This leads to a natural generalization of the concept of a straight line to that of a *geodesic*, where the geodesic between two points $u, v \in M$ is defined as a continuously differentiable curve of minimal length between them. These geodesics in Euclidean geometry are straight lines, and in spherical geometry they are arcs of great circles. We can similarly define the distance between two points, $d(x, y)$ as the length of a geodesic between them. In Euclidean space the relationship between a pair of nodes is defined along lines: the distance between the two nodes is the length of the line segment between them and forces between the two nodes act along the line through them. These notions of distance and forces can be extended to a Riemannian geometry by having these same relationships be defined in terms of the geodesics of the geometry, rather than in terms of Euclidean lines.

As Riemannian manifolds have a well-structured tangent space at every point, these tangent spaces can be used to generalize spring embedders to arbitrary Riemannian geometries. In particular, the tangent space is useful in dealing with the interaction between one point and several other points in non-Euclidean geometries. Consider three points x, y , and z in a Riemannian manifold M where there is an attractive force from x to y and z . As can be easily seen in the Euclidean case (but also true in general) the net force on x is not necessarily in the direction of y or z , and thus the natural motion of x is along neither the geodesic toward y , nor that toward z . Determining the direction in which x should move requires the notion of angle.

Since the tangent space at x , being an inner product space, has enough structure to define lengths and angles, we do the computations for calculating the forces on x in $T_x M$. In order

to do this, we define two functions for every point $x \in M$ as follows:

$$\begin{aligned}\tau_x &: M \rightarrow T_x M \\ \tau_x^{-1} &: T_x M \rightarrow M.\end{aligned}$$

These two functions map points in M to and from the tangent space of M at x , respectively. We require that τ_x and τ_x^{-1} satisfy the following constraints:

1. $\tau_x^{-1}(\tau_x(y)) = y$ for all $y \in M$
2. $\|\tau_x(y)\| = d(x, y)$
3. τ_x preserves angles about the origin

Using these functions it is now easy to define the way in which the nodes of a given graph $G = (V, E)$ interact with each other through forces. In the general framework for this algorithm each node is considered individually, and its new position is calculated based on the relative locations of the other nodes in the graph (repulsive forces) and on its adjacent edges (attractive forces). Then we obtain pseudo-code for a traditional Euclidean spring embedder and its corresponding non-Euclidean counterpart as follows:

```

generic_algorithm( $G$ )
while not done do
  foreach  $n \in G$  do
    position[ $n$ ] := force_directed_placement( $n, G$ )
  end
non_Euclidean_algorithm( $G$ )
while not done do
  foreach  $n \in G$  do
     $x$  := position[ $n$ ]
     $G'$  :=  $\tau_x(G)$ 
     $x'$  := force_directed_placement( $n, G'$ )
    position[ $n$ ] :=  $\tau_x^{-1}(x')$ 
  end
end

```

Bibliography

- [BBS97] Jürgen Branke, Frank Bucher, and Hartmut Schneck. A genetic algorithm for drawing undirected graphs. In *Proceedings of the 3rd Nordic Workshop on Genetic Algorithms and Their Applications*, pages 193–206, 1997.
- [BF96] I. Bruß and A. Frick. Fast interactive 3-D graph visualization. In F. J. Brandenburg, editor, *Proceedings of the 3rd Symposium on Graph Drawing (GD)*, volume 1027 of *Lecture Notes Computer Science*, pages 99–110. Springer-Verlag, 1996.
- [BH86] Josh Barnes and Piet Hut. A hierarchical $O(N \log N)$ force calculation algorithm. *Nature*, 324:446–449, December 1986.
- [Bra01] Ulrik Brandes. Drawing on physical analogies. In Michael Kaufmann and Dorothea Wagner, editors, *Drawing Graphs*, volume 2025 of *Lecture Notes in Computer Science*, pages 71–86. Springer-Verlag, 2001.
- [CLRS90] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. *Introduction to Algorithms*. MIT Press, Cambridge, MA, 1990.
- [Coh97] Jonathan D. Cohen. Drawing graphs to convey proximity: An incremental arrangement method. *ACM Transactions on Computer-Human Interaction*, 4(3):197–229, September 1997.
- [DETT99] Giuseppe Di Battista, Peter Eades, Roberto Tamassia, and Ioannis G. Tollis. *Graph Drawing: Algorithms for the Visualization of Graphs*. Prentice Hall, Englewood Cliffs, NJ, 1999.
- [DH96] Ron Davidson and David Harel. Drawing graphs nicely using simulated annealing. *ACM Transactions on Graphics*, 15(4):301–331, 1996.
- [Ead84] Peter Eades. A heuristic for graph drawing. *Congressus Numerantium*, 42:149–160, 1984.
- [EG95] Peter Eades and Patrick Garvan. Drawing stressed planar graphs in three dimensions. In *Proceedings of the 3rd Symposium on Graph Drawing*, pages 212–223, 1995.
- [Epp03] D. Eppstein. Hyperbolic geometry, Möbius transformations, and geometric optimization. In *MSRI Introductory Workshop on Discrete and Computational Geometry*, 2003.
- [FCW67] C. Fisk, D. Caskey, and L. West. Accel: Automated circuit card etching layout. In *Proceedings of the IEEE*, volume 55, pages 1971–1982, 1967.
- [FLM95] A. Frick, A. Ludwig, and H. Mehldau. A fast adaptive layout algorithm for undirected graphs. In R. Tamassia and I. G. Tollis, editors, *Proceedings of the 2nd Symposium on Graph Drawing (GD)*, volume 894 of *Lecture Notes in Computer Science*, pages 388–403. Springer-Verlag, 1995.
- [FR91] T. Fruchterman and E. Reingold. Graph drawing by force-directed placement. *Softw. – Pract. Exp.*, 21(11):1129–1164, 1991.
- [GGK04] P. Gajer, M. T. Goodrich, and S. G. Kobourov. A fast multi-dimensional algorithm for drawing large graphs. *Computational Geometry: Theory and Applications*, 29(1):3–18, 2004.
- [GGS03] C. Gotsman, X. Gu, and A. Sheffer. Fundamentals of spherical parameterization for 3D meshes. In *ACM Transactions on Graphics*, 22, pages 358–363, 2003.
- [Gre88] Leslie Greengard. *The Rapid Evolution of Potential Fields in Particle Systems*. M.I.T. Press, Cambridge, MA, 1988.

- [HH01] R. Hadany and D. Harel. A multi-scale algorithm for drawing graphs nicely. *Discrete Applied Mathematics*, 113(1):3–21, 2001.
- [HJ04] S. Hachul and M. Jünger. Drawing large graphs with a potential-field-based multilevel algorithm. In *Proceedings of the 12th Symposium on Graph Drawing (GD)*, volume 3383 of *Lecture Notes in Computer Science*, pages 285–295. Springer-Verlag, 2004.
- [HK02] David Harel and Yehuda Koren. A fast multi-scale method for drawing large graphs. *Journal of Graph Algorithms and Applications*, 6(3):179–2002, 2002.
- [Hoc96] *Approximation Algorithms for NP-Hard Problems*. PWS Publishing, 1996.
- [KGV83] S. Kirkpatrick, C. D. Gelatt, and M. P. Vecchi. Optimization by simulated annealing. *Science*, 220(4598):671–680, 1983.
- [KK89] T. Kamada and S. Kawai. An algorithm for drawing general undirected graphs. *Inform. Process. Lett.*, 31:7–15, 1989.
- [KMS91] Corey Kosak, Joe Marks, and Stuart Shieber. A parallel genetic algorithm for network-diagram layout. In *Proceedings of the 4 International Conference on Generic Algorithms*, pages 458–465, 1991.
- [KW05] S. G. Kobourov and K. Wampler. Non-Euclidean spring embedders. *IEEE Transactions on Visualization and Computer Graphics*, 11(6):757–767, 2005.
- [LRP95] John Lamping, Ramana Rao, and Peter Pirolli. A focus+context technique based on hyperbolic geometry for visualizing large hierarchies. In *Proceedings of Computer Human Interaction*, pages 401–408. ACM, 1995.
- [MB96] T. Munzner and P. Burchard. Visualizing the structure of the World Wide Web in 3D hyperbolic space. In David R. Nadeau and John L. Moreland, editors, *1995 Symposium on the Virtual Reality Modeling Language, VRML '95*, pages 33–38, 1996.
- [Mun97] Tamara Munzner. H3: Laying out large directed graphs in 3D hyperbolic space. In L. Lavagno and W. Reisig, editors, *Proceedings of IEEE Symposium on Information Visualization*, pages 2–10, 1997.
- [Mun98] T. Munzner. Drawing large graphs with h3viewer and site manager. In *Proceedings of the 6th Symposium on Graph Drawing*, pages 384–393, 1998.
- [OR01] J. Ontrup and H. Ritter. Hyperbolic self-organizing maps for semantic navigation. In *Advances in Neural Information Processing Systems 14*, pages 1417–1424, 2001.
- [Ost96] Diethelm Ironi Ostry. Some three-dimensional graph drawing algorithms. Master's thesis, University of Newcastle, Australia, 1996.
- [QB79] N. Quinn and M. Breur. A force directed component placement procedure for printed circuit boards. *IEEE Transactions on Circuits and Systems*, CAS-26(6):377–388, 1979.
- [QE00] Aaron Quigley and Peter Eades. FADE: graph drawing, clustering, and visual abstraction. In *Proceedings of the 8th Symposium on Graph Drawing (GD)*, volume 1984 of *Lecture Notes in Computer Science*, pages 197–210. Springer-Verlag, 2000.
- [Rit99] H. Ritter. Self-organizing maps on non-euclidean spaces. In S. Oja, E. & Kaski, editor, *Kohonen Maps*, pages 97–110. Elsevier, Amsterdam, 1999.
- [Tut63] William T. Tutte. How to draw a graph. *Proc. London Math. Society*, 13(52):743–768, 1963.
- [v-s99] *The Simple Genetic Algorithm: Foundations and Theory*. MIT Press, 1999.
- [vWN04] J. J. van Wijk and W. A. A. Nuij. A model for smooth viewing and navigation of large 2D information spaces. *IEEE Transactions on Visualization and Computer Graphics*, 10(4):447–458, 2004.

- [Wal03] C. Walshaw. A multilevel algorithm for force-directed graph drawing. *Journal of Graph Algorithms and Applications*, 7(3):253–285, 2003.

Index

chapter identifier, 1
chapter file, 1

environments, 2

figures, 2

labels, 2

main file, 1

packages, 1
pseudocode, 2